Uday Patel

uday.patel@sjsu.edu | linkedin.com/in/uday-patel-94a2a41ab | github.com/UdayPate | San Jose, CA

Education

San Jose State University — San Jose, CA
Master of Science in Applied Data Intelligence — *Expected Dec 2027*

- Coursework: Machine Learning, Deep Learning, Agentic AI & Distributed Systems, Data Mining, Data Visualization, Database Systems

University of California – Santa Cruz — Santa Cruz, CA
Bachelor of Science in Computer Science (GPA: 3.61) — *Jul 2021 – Jun 2025*

Experience

GTM Data Analyst Intern — Jun 2026 – Aug 2026
Everlaw — *Oakland, CA*

- Designed and shipped a first-of-its-kind post-sales benchmarking dashboard covering approximately 1,000 accounts (all-time) that lets Customer Success compare any account's renewal pacing against peers matched by sales segment, company size, and revenue.
- Built a custom dbt model unifying Salesforce, Gainsight, and existing dbt models with inconsistent account-level schemas into a single account-level view, resolving granularity mismatches across sources.
- Compared renewed vs. non-renewed AI-subscription accounts to surface what separated the groups, and flagged low-data accounts in the dashboard so CS owners can see where tracking is thin.
- Validated the approach with CS and GTM stakeholders throughout; the dashboard moved to staging for company-wide rollout.
- Built 4 Tableau dashboards on AI subscription growth, active-usage rates, and EverlawAI Deep Dive adoption, supporting GTM, Marketing, and Customer Success teams; insights presented by the CMO and CEO at the quarterly all-hands.
- Rebuilt the 4 dashboards in Looker using LookML (parameters, derived tables, rolling-window calculations) and authored a Tableau-vs-Looker comparison the team used to evaluate its long-term reporting stack.
- Debugged and rewrote production dbt models attributing Customer Success and Account Executive ownership at key account milestones, fixing broken sales-attribution logic and improving downstream reporting accuracy.

Apprentice – Data Analytics — Jul 2024 – Sep 2024
GroovyWeb

- Analyzed an A/B test of client landing-page designs, comparing conversion rates across variants and recommending the higher-performing design to guide website optimization.
- Built Tableau dashboards from client business data to support client-facing presentations and identify performance improvement opportunities.
- Wrote SQL queries to extract, join, and analyze data across multiple business systems for recurring client reporting and decision support.
- Collaborated with cross-functional teammates in an Agile/Scrum environment to translate client business needs into data-driven dashboards and reporting solutions.

Projects

Data Lens: Automated Dashboard Recommendation Platform | ***PySpark, AWS, Airflow, dbt, LLM Agents*** Aug 2026

- Built an Airflow-orchestrated pipeline that ingests arbitrary CSV/JSON/Parquet datasets into S3, infers schemas, and transforms them into analysis-ready tables with Glue/PySpark and dbt, with automated data-quality tests; processes datasets up to 1M rows.
- Engineered a distributed profiling engine in PySpark computing column-level statistics and inferential tests (chi-square, correlations with p-values, changepoint detection) across 15 datasets in 3 domains, cutting profiling time from 8 minutes to 90 seconds.
- Designed an agent layer that recommends dashboards and charts grounded only in the computed statistical evidence, evaluated with a cross-domain harness scoring precision, recall, and faithfulness.

Airbnb Data Engineering Pipeline & Graph Recommender | ***Snowflake, dbt, Python, AWS S3, Neo4j*** Mar 2026

- Staged source files in AWS S3 and loaded them into Snowflake, then built a Medallion pipeline (Bronze/Silver/Gold) with dbt: 11 incremental models over 5,000 booking records with SCD Type 2 historical tracking and a star-schema Gold layer.
- Performed feature engineering in the Silver layer (host tier segmentation, capacity categorization) via Jinja-templated, meta-driven SQL to standardize metric definitions across datasets.
- Exported the Gold layer to Neo4j AuraDB (1,189 nodes, 5,500 relationships) and built a Flask REST API with 6 endpoints serving item-based collaborative filtering recommendations via Cypher.

Apple Sales & Warranty Analytics Platform | ***PostgreSQL, SQL, Streamlit, Plotly*** Jan 2026

- Architected a normalized 5-table PostgreSQL database modeling product sales and warranty claims across 1M+ records, ensuring data integrity and scalability.
- Engineered 16+ analytical queries using CTEs, window functions, and multi-dimensional aggregations; improved database performance by 85–97% through strategic indexing.
- Implemented linear regression in pure SQL using closed-form equations for demand forecasting, validated against scikit-learn.
- Developed an interactive multi-page Streamlit dashboard with real-time SQL querying and Plotly visualizations for executive-level reporting.

AI Cricket Intelligence Platform | ***Python, PyTorch, Transformers, RAG, PostgreSQL, FastAPI*** Sep 2025

- Built an end-to-end AI platform predicting IPL player runs, wickets, and win probabilities, training LSTM and Transformer models on sequential historical data, player form, opponent strength, and venue context; Transformer reached 15.7 runs MAE vs. 18.4 for LSTM.
- Engineered a RAG pipeline with embeddings and vector search to retrieve player, team, venue, and historical-match statistics, grounding predictions in evidence-backed data.
- Integrated an LLM to generate natural-language explanations and built an interactive Streamlit dashboard for predictions, player comparisons, and AI-generated match analysis.

Cravely: Recipe Discovery Web App | ***Py4Web, Vue.js, HTML, CSS, PostgreSQL*** Mar 2025

- Led the design and implementation of a full-stack recipe platform, defining features, user flows, and backend architecture to improve recipe discovery.
- Engineered RESTful APIs for recipe and ingredient CRUD operations, authentication/authorization, and image handling, integrating external data from the TheMealDB API.

Technical Skills

Languages: Python, SQL (CTEs, window functions, query optimization), JavaScript, C/C++, HTML/CSS
Data Engineering: Snowflake, dbt (incremental models, SCD Type 2, Jinja), Airflow, PySpark, AWS (S3, Glue, Athena), ETL/ELT, Medallion Architecture, Data Modeling (star schema)
Databases: PostgreSQL, Neo4j
BI & Analytics: Tableau, Looker (LookML), Power BI, Excel, Plotly, Streamlit, A/B Testing, EDA, Metrics Design
ML & AI: PyTorch, TensorFlow, scikit-learn, Transformers, LSTM, RAG, LLM Integration & Agents, pandas, NumPy
Tools: Git/GitHub, Salesforce, Gainsight, Flask, FastAPI, Jupyter Notebook, Jira, Cursor, GitHub Copilot, Claude